

#前言

这篇文章将用实战来表明如何用 `gdb` 从 `coredump` 恢复现场并游走在程序的堆栈之间。

#问题

双11之前, 为了使用我的工具集 `systemtap-toolkit` 进行生产环境的活体分析, 需要给 `php` 开启调试符号, 开启方式比较简单, 如下所示,

```
./configure --enable-debug ....
```

重新打包发布测试, 按理论应该不会有问题, 毕竟只是开启调试符号, 并没有修改 codebase。预发测试没问题, ok, 生产上线部分流量看看。WTF `PHP-FPM` 经常的报 `coredump`。

因缺思婷, 开调试符号会造成 `coredump`, 还是第一次见。这个现象成功地吸引了我的注意力。

#载入coredump

扔出调试工具 `gdb` 查查原因。

先载入 `coredump` 恢复现场。

```
# gdb -c /tmp/core-500-php-fpm-29093-1478780406 $(php-fpm)
Excess command line arguments ignored. (/opt/isys/sbin/php-fpm ...)
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-90.el6)
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>...

Core was generated by php-fpm: pool www.
Program terminated with signal 11, Segmentation fault.
#0  0x0000000000a6e6a5 in _zval_dtor_func (zvalue=0x7f676a8a93c8,
    __zend_filename=0x1090db8 "/data/users/app/rpmbuild/BUILD/php-5.6.16/Zend/zend_execute.h"
36      CHECK_ZVAL_STRING_REL(zvalue);
```

#调用栈

看看具体的调用栈。

```
(gdb) bt
#0  0x000000000a6e6a5 in _zval_dtor_func (zvalue=0x7f676a8a93c8,
    __zend_filename=0x1090db8 "/data/users/app/rpmbuild/BUILD/php-5.6.16/Zend/zend_execute.h
#1  0x000000000a59abe in _zval_dtor (zvalue=0x7f676a8a93c8,
    __zend_filename=0x1090db8 "/data/users/app/rpmbuild/BUILD/php-5.6.16/Zend/zend_execute.h
#2  0x000000000a59b8a in i_zval_ptr_dtor (zval_ptr=0x7f676a8a93c8,
    __zend_filename=0x1092fb0 "/data/users/app/rpmbuild/BUILD/php-5.6.16/Zend/zend_variables
#3  0x000000000a5add2 in _zval_ptr_dtor (zval_ptr=0x7f676a8a96b0,
    __zend_filename=0x1092fb0 "/data/users/app/rpmbuild/BUILD/php-5.6.16/Zend/zend_variables
#4  0x000000000a6ec1c in _zval_ptr_dtor_wrapper (zval_ptr=0x7f676a8a96b0)
    at /usr/src/debug/php-5.6.16/Zend/zend_variables.c:188
#5  0x000000000a82bd1 in i_zend_hash_bucket_delete (ht=0x142c328, p=0x7f676a8a9698)
    at /usr/src/debug/php-5.6.16/Zend/zend_hash.c:182
#6  0x000000000a82ca8 in zend_hash_bucket_delete (ht=0x142c328, p=0x7f676a8a9698)
    at /usr/src/debug/php-5.6.16/Zend/zend_hash.c:192
#7  0x000000000a849d9 in zend_hash_graceful_reverse_destroy (ht=0x142c328)
    at /usr/src/debug/php-5.6.16/Zend/zend_hash.c:613
---Type <return> to continue, or q <return> to quit---
#8  0x000000000a5a5f8 in shutdown_executor ()
    at /usr/src/debug/php-5.6.16/Zend/zend_execute_API.c:244
#9  0x000000000a70e61 in zend_deactivate () at /usr/src/debug/php-5.6.16/Zend/zend.c:960
#10 0x0000000009d6b5d in php_request_shutdown (dummy=0x0)
    at /usr/src/debug/php-5.6.16/main/main.c:1883
#11 0x000000000b3633b in main (argc=4, argv=0x7ffcc1123638)
    at /usr/src/debug/php-5.6.16/sapi/fpm/fpm/fpm_main.c:1992
```

看看产生 `coredump` 的代码, 注意第36行:)

```
(gdb) list _zval_dtor_func
27 #include "zend_constants.h"
28 #include "zend_list.h"
29
30
31 ZEND_API void _zval_dtor_func(zval *zvalue ZEND_FILE_LINE_DC)
32 {
33     switch (Z_TYPE_P(zvalue) & IS_CONSTANT_TYPE_MASK) {
34         case IS_STRING:
35         case IS_CONSTANT:
36             CHECK_ZVAL_STRING_REL(zvalue);
37             str_efree_rel(zvalue->value.str.val);
38             break;
39         case IS_ARRAY: {
40             TSRMLS_FETCH();
41
42             if (zvalue->value.ht && (zvalue->value.ht != &EG(symbol_table))) {
43                 /* break possible cycles */
44                 Z_TYPE_P(zvalue) = IS_NULL;
45                 zend_hash_destroy(zvalue->value.ht);
46                 FREE_HASHTABLE(zvalue->value.ht);
```

`CHECK_ZVAL_STRING_REL(zvalue);` 看起来是个宏, 翻出 php-src 的源码查查 `CHECK_ZVAL_STRING_REL` 的定义.

#分析PHP源码

```
git clone https://github.com/php/php-src.git
git checkout php-5.6.16
```

用 [cloc](#) 统计下 [php](#) 的代码量:)

```
> cloc .
1000 files
14600 files
14700 files

15000 files
15100 files
17131 text files.
16869 unique files.
14491 files ignored.
```

2s

<http://cloc.sourceforge.net> v 1.62 T=121.67 s (21.2 files/s, 11662.1 lines/s)

Language	files	blank	comment	code
C	996	102590	133526	841438
C/C++ Header	749	20249	29558	184948
PHP	448	4740	7718	23523
m4	110	2421	763	18026
Bourne Shell	35	1356	1863	12522
XML	86	177	117	6263
C++	20	1290	738	6198
Javascript	3	537	219	2464
lex	3	546	343	2317
Windows Module Definition	7	1	6	1918
yacc	3	293	108	1564
Expect	10	127	17	1515
Pascal	29	260	1471	1058
make	13	149	37	716
awk	10	87	53	541
HTML	7	37	0	368
Perl	1	81	28	316
Windows Resource File	7	265	81	302
XSD	9	33	0	265
XSLT	12	16	24	160
D	4	12	38	149
DOS Batch	7	12	4	69
YAML	2	8	3	59
DTD	5	11	0	46
SQL	1	4	0	40
NAnt script	1	8	0	25
Windows Message File	1	4	0	24
Ant	1	5	0	10
SUM:	2580	135319	176715	1106844

看起来很大, 估计用[ag](#)搜会搜出一坨:(。ok, 拿 [cscope](#) 来查比较精确。

```
cscope -Rbq
cscope -dq
```

Global definition: CHECK_ZVAL_STRING_REL

	File	Line
0	zend_API.h	539 <i>#define CHECK_ZVAL_STRING_REL(z) \</i>
1	zend_API.h	543 <i>#define CHECK_ZVAL_STRING_REL(z)</i>

Find this C symbol:
Find this global definition: CHECK_ZVAL_STRING_REL
Find functions called by this
Find functions calling this function:
Find this text string:
Change this text string:
Find this egrep pattern:
Find this file:
Find files *#including this file*:
Find assignments to this symbol:

看看具体的 CHECK_ZVAL_STRING_REL 定义

```
539 #define CHECK_ZVAL_STRING_REL(z) \
540 >   if (Z_STRVAL_P(z)[ Z_STRLEN_P(z) ] != '\0') { zend_error(E_WARNING, "String is not ze
```

看起来 PHP 宏用的飞起，呵呵。看看 Z_STRVAL_P 的定义(方法如上)

```
460 #define Z_STRVAL_P(zval_p)  Z_STRVAL(*zval_p)
461 #define Z_STRLEN_P(zval_p)  Z_STRLEN(*zval_p)
444 #define Z_STRVAL(zval)      (zval).value.str.val
445 #define Z_STRLEN(zval)      (zval).value.str.len
```

所以最后的表达式就是 `(*zvalue).value.str.val[(zvalue).value.str.len]`，再看看这几个变量的值。

```
(gdb) p (*zvalue).value.str.val
$3 = 0x0
(gdb) p (*zvalue).value.str.len
$4 = 0
```

显然，程序访问了非法的内存地址0x00.(有兴趣的自己查看程序布局`cat /proc/xx/maps`)。造成 [coredump](#) 的代码找到了，现在得想办法找出复现条件，由于是在发生 [coredump](#) 三天之后才开始调试的，之前的现场只有一份 [coredump](#)。

#复现现场

曾经用 [Node.js](#) 实现过 [fastcgi](#) 协议，我猜测代码里应该有保存请求信息的对象，仔细看调用栈，看起来入口比较像。

```
(gdb) frame 11
#11 0x0000000000b3633b in main (argc=4, argv=0x7ffcc1123638)
    at /usr/src/debug/php-5.6.16/sapi/fpm/fpm/fpm_main.c:1992
1992      php_request_shutdown((void *) 0);
```

看看具体的代码是啥

```
(gdb) list /usr/src/debug/php-5.6.16/sapi/fpm/fpm/fpm_main.c:1992
1987      fpm_log_write(NULL TSRMLS_CC);
1988
1989      STR_FREE(SG(request_info).path_translated);
1990      SG(request_info).path_translated = NULL;
1991
1992      php_request_shutdown((void *) 0);
1993
1994      requests++;
1995      if (max_requests && (requests == max_requests)) {
1996          fcgi_finish_request(&request, 1);
```

SG(request_info) 看起来比较像:)，看看 SG 是啥

```
147 # define SG(v) (sapi_globals.v)
```

从命名方式来看像是全局变量，打印出值看看:) (由于隐私问题，部分信息被修改，但不影响)

```
(gdb) p sapi_globals.request_info
$1 = {request_method = 0x7f676a876560 "POST",
      query_string = 0x7f676a876430 "qqqq=123",
      cookie_data = 0x7f676a878250 "a=1;b=2"... , content_length = 0, path_translated = 0x0,
      request_uri = 0x7f676a876880 "/yyyy", request_body = 0x7f676a8a8e98,
      content_type = 0x7f676a876670 "", headers_only = 0 '\000', no_headers = 0 '\000',
      headers_read = 0 '\000', post_entry = 0x0, content_type_dup = 0x7f676a8a8dc0 "",
      auth_user = 0x0, auth_password = 0x0, auth_digest = 0x0, argv0 = 0x0, current_user = 0x0,
      current_user_length = 0, argc = 0, argv = 0x0, proto_num = 1000}
```

Good，试试复现构造这个请求。

```
curl -X POST -H "Cookie: a=1;b=2" -H "Content-Length: 0" http://127.0.0.1/yyy?qqqq=123
```

果然，POST 一个 request body 为空的请求比较少见，竟然造成 [coredump](#) :(

到目前为止，复现方法和 [coredump](#) 原因都已经找到，再继续挖掘调用栈看看，从外到里一层层看，在 frame 8 发现新线索:)

#新线索

```
(gdb) frame 8
#8 0x0000000000a5a5f8 in shutdown_executor ()
    at /usr/src/debug/php-5.6.16/Zend/zend_execute_API.c:244
244      zend_hash_graceful_reverse_destroy(&EG(symbol_table));
```

看语义是在销毁 EG(symbol_table) 这个 hash 造成的，看看具体 hash 里是什么:)

```
47 # define EG(v) (executor_globals.v)
```

看起来又是一个全局变量，打印出看看：)

```
(gdb) p executor_globals.symbol_table
$2 = {nTableSize = 64, nTableMask = 63, nNumOfElements = 0, nNextFreeElement = 0,
      pInternalPointer = 0x0, pListHead = 0x0, pListTail = 0x0, arBuckets = 0x7f676a8a9440,
      pDestructor = 0xa6ebf8 <_zval_ptr_dtor_wrapper>, persistent = 0 '\000',
      nApplyCount = 0 '\000', bApplyProtection = 1 '\001', inconsistent = 0}
```

结构比较复杂，从成员的命名方式可以知道应该是个链表。看看具体的 `symbol_table` 类型

```
168 struct _zend_executor_globals {
169 >   zval **return_value_ptr_ptr;
170
171 >   zval uninitialized_zval;
172 >   zval *uninitialized_zval_ptr;
173
174 >   zval error_zval;
175 >   zval *error_zval_ptr;
176
177 >   /* symbol table cache */
178 >   HashTable *symtable_cache[SYMTABLE_CACHE_SIZE];
179 >   HashTable **symtable_cache_limit;
180 >   HashTable **symtable_cache_ptr;
181
182 >   zend_op **opline_ptr;
183
184 >   HashTable *active_symbol_table;
185 >   HashTable symbol_table;>>   /* main symbol table */
```

果然跟之前说的一样，是个 `HashTable`。既然是 `HashTable`，应该有记录 hash 对应的 key 值，看看 `HashTable` 的定义

```
67 typedef struct _hashtable {
68 >   uint nTableSize;
69 >   uint nTableMask;
70 >   uint nNumOfElements;
71 >   ulong nNextFreeElement;
72 >   Bucket *pInternalPointer;>   /* Used for element traversal */
73 >   Bucket *pListHead;
74 >   Bucket *pListTail;
75 >   Bucket **arBuckets;
76 >   dtor_func_t pDestructor;
77 >   zend_bool persistent;
78 >   unsigned char nApplyCount;
79 >   zend_bool bApplyProtection;
80 #if ZEND_DEBUG
81 >   int inconsistent;
82 #endif
83 } HashTable;
```

看 `pListHead` 和 `pListTail`，看样子所有的变量都在这个链表连起来的桶上，看看 `Bucket` 的定义

```

55 typedef struct bucket {
56 >   ulong h;>   >   >   >   >   /* Used for numeric indexing */
57 >   uint nKeyLength;
58 >   void *pData;
59 >   void *pDataPtr;
60 >   struct bucket *pListNext;
61 >   struct bucket *pListLast;
62 >   struct bucket *pNext;
63 >   struct bucket *pLast;
64 >   const char *arKey;
65 } Bucket;

```

arKey 应该就是我们要找的 key。ok，调用栈上应该有具体的某个Bucket，往里继续找着:)

```

(gdb) frame 7
#7  0x0000000000a849d9 in zend_hash_graceful_reverse_destroy (ht=0x142c328)
    at /usr/src/debug/php-5.6.16/Zend/zend_hash.c:613
613      zend_hash_bucket_delete(ht, ht->pListTail);

(gdb) list /usr/src/debug/php-5.6.16/Zend/zend_hash.c:613
608 ZEND_API void zend_hash_graceful_reverse_destroy(HashTable *ht)
609 {
610     IS_CONSISTENT(ht);
611
612     while (ht->pListTail != NULL) {
613         zend_hash_bucket_delete(ht, ht->pListTail);
614     }
615
616     if (ht->nTableMask) {
617         pefree(ht->arBuckets, ht->persistent);

```

看代码是在遍历 ht->pListTail 的某一个值导致的，继续往里看看发生 [coredump](#) 时的 ht->pListTail 值是多少

```

#6  0x0000000000a82ca8 in zend_hash_bucket_delete (ht=0x142c328, p=0x7f676a8a9698)
    at /usr/src/debug/php-5.6.16/Zend/zend_hash.c:192
192     i_zend_hash_bucket_delete(ht, p);

```

打印出来看看

```

(gdb) p p
$4 = (Bucket *) 0x7f676a8a9698
(gdb) p *p
$5 = {h = 12508711195225833452, nKeyLength = 19, pData = 0x7f676a8a96b0,
      pDataPtr = 0x7f676a8a93c8, pListNext = 0x0, pListLast = 0x0, pNext = 0x0, pLast = 0x0,
      arKey = 0x7f676a8a96e0 "HTTP_RAW_POST_DATA"}

```

奈斯，看样子是在销毁 "HTTP_RAW_POST_DATA" 变量导致的 [coredump](#)，看看 HTTP_RAW_POST_DATA 那里定义:)

Text string: HTTP_RAW_POST_DATA

```
File      Line
0 php_var.h      157 (name_len == sizeof("HTTP_RAW_POST_DATA") - 1 && !memcmp(name,
    "HTTP_RAW_POST_DATA", sizeof("HTTP_RAW_POST_DATA") - 1)) ||
1 php_content_types.c 69 "HTTP_RAW_POST_DATA truncated from %lu to %d bytes",
2 php_content_types.c 73 SET_VAR_STRINGL("HTTP_RAW_POST_DATA", data,

);
3 php_content_types.c 76 "Automatically populating $HTTP_RAW_POST_DATA is deprecated and "
```

Find this C symbol:

Find this global definition:

Find functions called by this function:

Find functions calling this function:

Find this text string:

Change this text string:

Find this egrep pattern:

Find this file:

Find files *#including this file:*

Find assignments to this symbol:

```
59 > > if (populate_raw_post_data(TSRMLS_C)) {
60 > > > size_t length;
61 > > > char *data = NULL;
62
63 > > > php_stream_rewind(SG(request_info).request_body);
64 > > > length = php_stream_copy_to_mem(SG(request_info).request_body, &data, PHP_ST
65 > > > php_stream_rewind(SG(request_info).request_body);
66
67 > > > if (length > INT_MAX) {
68 > > > > sapi_module.sapi_error(E_WARNING,
69 > > > > > "HTTP_RAW_POST_DATA truncated from %lu to %d bytes",
70 > > > > > (unsigned long) length, INT_MAX);
71 > > > > length = INT_MAX;
72 > > > }
73 > > > SET_VAR_STRINGL("HTTP_RAW_POST_DATA", data, length);
```

注意第73行 `data` 和 `length`，跟我们之前 `(*zvalue).value.str.val[(zvalue).value.str.len]` 的相对应，当我们发起一个 POST Request Body为空的请求时，`data`将为空(这里可以直接调试php解释器验证)，ok，真正的根源找到了看看怎么修复。直觉告诉我 `(*zvalue).value.str.val[(zvalue).value.str.len]`，之前应该加个判断，不过在不熟悉源码的情况下，还是改 `HTTP_RAW_POST_DATA` 比较安全：) 当 `data` 为 `NULL` 并且 `length` 为 0 时，只要给 `data` 赋予一个合法的地址，并保证 `Z_STRVAL_P(z)[Z_STRLEN_P(z)]` 为 `'\0'` 即可，最粗暴的方式直接用 `data = ""`；试试。

重新编译打包依然 `coredump`，不过这次 `coredump` 的原因不一样。

```
#0 0x0000000000a34a18 in zend_mm_check_ptr (heap=0x30bf2d0, ptr=0x107db83, silent=1,
    __zend_filename=0x1090558 "/root/rpmbuild/SOURCES/php-5.6.16/Zend/zend_execute.h",
    __zend_lineno=79,
    __zend_orig_filename=0x1092700 "/root/rpmbuild/SOURCES/php-5.6.16/Zend/zend_variables.c"
1384 if (p->info._size != ZEND_MM_NEXT_BLOCK(p)->info._prev) {
```


看起来不是堆上的地址造成的 [coredump](#) , [PHP](#) 应该有一套自己的内存管理机制。没办法了, 只能找找对应的 API, 按经验试试搜索 `empty` 和 `str`。

Text string: empty

File	Line
0 zend.c	326 <code>if (Z_STRLEN_P(expr) == 0) { /* optimize away empty strings */</code>
1 zend.c	525 <code>zend_hash_copy(compiler_globals->auto_globals, global_auto_globals_table, NULL, NULL, sizeof(zend_auto_global) /* empty element */);</code>
2 zend.h	682 <code>#define STR_EMPTY_ALLOC() CG(interned_empty_string)? CG(interned_empty_string) : estrndup("", sizeof(""))-1;</code>

* Lines 1-4 of 2347, 2344 more - press the space bar to display more *

Find this C symbol:

Find this global definition:

Find functions called by this function:

Find functions calling this function:

Find this text string:

Change this text string:

Find this `egrep` pattern:

Find this file:

Find files *#including this file:*

Find assignments to this symbol:

Good Luck, 看 `STR_EMPTY_ALLOC()` 应该是, 加了以下代码直接重新编译测跑测试通过。

```
if (data == NULL && length == 0) {
    data = STR_EMPTY_ALLOC();
}
```

#总结

如果之前熟悉 [php](#) 源码, 拿到调用栈直接打印 `hash` 的 `key`, 估计十几分钟内就完事, 最后不得不花了几个小时才搞定。一个悲伤的事是当我准备顺手提交 PR 给 [PHP](#) 官方时, 搜了下 [Changelog](#), 发现早有人提交并修复了这个问题, 如果你的情况比较紧急, 建议还是看看新版本是否修复吧, 至于我只是为了好玩哈哈。